

Week 3: R Functions

Sarah Tareen

2023-04-21

This week we are introducing **R functions** and how to write our own functions.

Questions to answer:

Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped. Your final function should be adequately explained with code comments and be able to work on an example class gradebook such as this one in CSV format: “<https://tinyurl.com/gradeinput>” [3pts]

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)

student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)

student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

Follow the guidelines from class

- Write a working snippet of code that solves a simple problem

```
# Straightforward mean()
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)

mean(student1)
```

```
## [1] 98.75
```

But... We need to drop the lowest score. First we need to identify the lowest score.

```
# Which element of the vector is the lowest?
which.min(student1)
```

```
## [1] 8
```

What I want is to now drop (ie. exclude) this lowest score from my `mean()` calculation.

```
# This will return everything but the eighth element of the vector.
student1[-8]
```

```
## [1] 100 100 100 100 100 100 100
```

Now we can use the answer from `which.min()` to return all other elements of the vector.

```
#This is our first working snippet!  
mean( student1[-which.min(student1)] )
```

```
## [1] 100
```

What about the other example students? Will this work for them?

We could try using the `na.rm = TRUE` argument for `mean` but this is pants! not a good approach i.e. unfair.

```
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)  
mean(student2, na.rm=TRUE)
```

```
## [1] 91
```

```
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)  
mean(student3, na.rm=TRUE)
```

```
## [1] 90
```

Another approach is to mask (i.e. replace) all NA values with zero

First we need to find the NA elements of the vector. How do we do this?

```
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)  
x <- student2  
  
is.na(x)
```

```
## [1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

```
which( is.na(x) )
```

```
## [1] 2
```

Now we have identified the NA elements, we want to “mask” them. Replace them with zero?

```
# This does not quite get us there  
mean(x[-which(is.na(x))])
```

```
## [1] 91
```

Instead we will make the NA elements zero

```
#Cool, this is useful  
x[is.na(x)] <- 0  
x
```

```
## [1] 100 0 90 90 90 90 97 80
```

```
mean(x)
```

```
## [1] 79.625
```

Recall we should drop the lowest score now..

```
x[is.na(x)] <- 0  
mean( x[-which.min(x)] )
```

```
## [1] 91
```

Now we are essentially there with our working snippet!

```
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)  
x <- student3  
x[is.na(x)] <- 0  
mean( x[-which.min(x)] )
```

```
## [1] 12.85714
```

Now we can make our function

Take the snippet and turn it into a function. Every function has 3 parts:

- A name, in our case 'grade()'
- Input arguments, a vector of student scores
- The body i.e. our working snippet of code

Using RStudio I will select code then do **Code > Extract Function**

```
grade <- function(x) {  
  x[is.na(x)] <- 0  
  mean( x[-which.min(x)] )  
}
```

```
grade(student1)
```

```
## [1] 100
```

```
grade(student2)
```

```
## [1] 91
```

```
grade(student3)
```

```
## [1] 12.85714
```

This looks great! We now need to add comments to explain this to our future selves and other who want to use this function. You can select code from bottom up and select **Code > Roxygen skeleton**.

```

#' Calculate the average score for a vector of student scores dropping the lowest score.
#' Missing values will be treated as zero.
#'
#' @param x A numeric vector of homework scores
#'
#' @return Average score
#' @export
#'
#' @examples
#' student <- c(100, NA, 90, 97)
#' grade(student)
#'
grade <- function(x) {
  # mask NA with zero
  # Treat missing values as zero
  x[is.na(x)] <- 0
  # Exclude the lowest score from mean
  mean( x[-which.min(x)] )
}

```

Now finally we can use our function on our “real” whole class data from this CSV format file: “https://tinyurl.com/gradeinput”

```

url <- "https://tinyurl.com/gradeinput"
gradebook <- read.csv(url, row.names = 1)

```

```

# Use this to apply the function to the rows of the data table
apply(gradebook, 1, grade)

```

```

## student-1 student-2 student-3 student-4 student-5 student-6 student-7
##      91.75      82.50      84.25      84.25      88.25      89.00      94.00
## student-8 student-9 student-10 student-11 student-12 student-13 student-14
##      93.75      87.75      79.00      86.00      91.75      92.25      87.75
## student-15 student-16 student-17 student-18 student-19 student-20
##      78.75      89.50      88.00      94.50      82.75      82.75

```

Q2. Using your grade() function and the supplied gradebook, Who is the top scoring student overall in the gradebook? [3pts]

To answer this we can run the apply() function and save the results.

```

results <- apply(gradebook, 1, grade)

sort(results, decreasing = TRUE)

```

```

## student-18 student-7 student-8 student-13 student-1 student-12 student-16
##      94.50      94.00      93.75      92.25      91.75      91.75      89.50
## student-6 student-5 student-17 student-9 student-14 student-11 student-3
##      89.00      88.25      88.00      87.75      87.75      86.00      84.25
## student-4 student-19 student-20 student-2 student-10 student-15
##      84.25      82.75      82.75      82.50      79.00      78.75

```

```
which.max(results)
```

```
## student-18  
##      18
```

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall)? [2pts]

```
gradebook
```

```
##      hw1 hw2 hw3 hw4 hw5  
## student-1 100 73 100 88 79  
## student-2 85 64 78 89 78  
## student-3 83 69 77 100 77  
## student-4 88 NA 73 100 76  
## student-5 88 100 75 86 79  
## student-6 89 78 100 89 77  
## student-7 89 100 74 87 100  
## student-8 89 100 76 86 100  
## student-9 86 100 77 88 77  
## student-10 89 72 79 NA 76  
## student-11 82 66 78 84 100  
## student-12 100 70 75 92 100  
## student-13 89 100 76 100 80  
## student-14 85 100 77 89 76  
## student-15 85 65 76 89 NA  
## student-16 92 100 74 89 77  
## student-17 88 63 100 86 78  
## student-18 91 NA 100 87 100  
## student-19 91 68 75 86 79  
## student-20 91 68 76 88 76
```


```
ave.scores <- apply(gradebook, 2, mean, na.rm=TRUE)  
which.min(ave.scores)
```

```
## hw3  
##    3
```

```
med.scores <- apply(gradebook, 2, median, na.rm=TRUE)  
which.min(med.scores)
```

```
## hw2  
##    2
```

```
#line shows median  
boxplot(gradebook)
```



Since the median is less sensitive to outliers we should use it as an average. Therefore, HW 2 got the most lowest scores overall.

Q4 From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)? [1pt]

Are the final results (i.e. average score for each student) correlated with the results (i.e. scores) for individual homeworks - the gradebook columns.

```
masked.gradebook <- gradebook
masked.gradebook[is.na(masked.gradebook)] <- 0

# And look at correlation
cor(results, masked.gradebook$hw1)
```

```
## [1] 0.4250204
```

```
apply(masked.gradebook, 2, cor, x=results)
```

```
##      hw1      hw2      hw3      hw4      hw5
## 0.4250204 0.1767780 0.3042561 0.3810884 0.6325982
```

HW 5 has the highest correlation value with the average scores and is most predictive of overall score.

Q5 Make sure you save your Quarto document and can click the “Render” (or Rmarkdown” Knit”) button to generate a PDF format report without errors. Finally, submit your PDF to gradescope. [1pt]

Knit the document to make the PDF